Is the Golden Age of Sexist Commentary Over? A Gendered Analysis of Curling Commentary at the 2026 Milan-Cortina Winter Olympic Games.

Victor Le Lay
ENSAE Paris
victor.lelay@ensae.fr

Abstract

Gender bias in sports media coverage has been extensively documented, yet live broadcast commentary remains understudied. This collective work investigates the prevalence of latent sexist bias in French-language curling commentary during the 2026 Milan-Cortina Winter Olympic Games. The corpus used includes 18 events (16.5 hours) retrieved from France Télévisions. Audio files were transcribed and annotated using Whisper, pyannote.audio, and InaSpeechSegmenter. We yield 11,909 labelled sentences with speaker and athlete gender metadata. We apply BERTopic with CamemBERT embeddings to identify thematic patterns across commentary segments. In complement, we use a cosine similarity analysis to test whether theoretically grounded terms are used in different semantic contexts depending on gender. Our results suggest that athlete gender generates limited thematic differentiation, consistently with a broader trend toward reduced explicit bias in sports commentary. However, commentator gender tends to produce slightly more substantial and numerous topic deviations, suggesting that latent biases, if they persist, are more concentrated on the journalists' side.

1 Introduction

In the late 2010s, the French public TV group implemented a policy aiming to yield gendered parity on screen. Executives claim the underlying ambition of fighting gender stereotypes and discrimination in sports. Such policies have a well-known effect on women's representation, especially among executive positions (De Acutis et al., 2024). However, their impact on implicit gender bias on live television is still under-studied and hard to measure. This paper's contribution relies on filling this specific gap.

Regarding gender biases in sport journalism, the quantitative literature mostly focuses on written press and post-event interviews rather than live commentary. Many methods are employed but when it comes to topic modeling, LDA remains the most used model for gender biases in sport journalism, as in Rao and Taboada, 2021. As far as our knowledge goes, there have been no implementation of BERTopic on the matter, even though it outperforms most of the other topic modeling tools (Mutsaddi et al., 2025). Though LLM-enhanced benchmarks are becoming more popular and can outperform a BERTopic, we chose to use tools for which we had 100% control over. Cosine distance, which we use for robustness analysis, is being more and more challenged (Zhou et al., 2022) but remains a relevant and simple measure of semantic similarity.

Since the 2000s, sexism became less blatant (Messner et al., 1993) and gradually more implicit (Cooky et al., 2015). Live commentary constitutes a promising field to unravel this trend by offering a large enough surface to potentially highlight latent bias, but also by enabling potential sexist "slips"

more than written articles, for example. If we expect to observe no blatant sexist interventions, we could either highlight the absence of latent sexism as a conclusion of an ongoing trend, or observe implicit sexist artifacts.

The lack of quantitative works on live sport commentary and the dated nature of qualitative works — especially regarding Olympic Games (Billings and Eastman, 2003; Capranica and Aversa, 2002) — led us to advocate for a necessary update on this matter. Furthermore, the political and social context pushes for feminist resistance to sexist biases, and a quantitative empirical work would thus contribute to fulfilling a social need. Finally, the exclusion of a Eurosport (UK) journalist during the 2024 Olympic Games following a sexist comment on live television[1] further illustrates the relevance of this inquiry.

I chose to apply topic modelling analysis to live commentary of curling events during the 2026 Winter Olympic Games in order to measure potential gendered biases. I test two hypotheses: topic distributions vary depending on the gender of both athletes (H1) and journalists (H2). This analytical framework constitutes, to my knowledge, the first use of NLP techniques to study French live sport commentary of the Olympic Games.

2 Data

This section aims to describe data and how there were gathered and pre-processed accordingly in order to conduct our analytic strategy.

2.1 Corpus collection

The corpus contains all curling events broadcast during the 2026 Milan-Cortina Winter Olympic Games. Lucas and Quentin cover other sports in their own reports. The files were retrieved from the France Télévisions streaming platform and therefore consist of French-language commentary. We extracted video files (.mp4) and converted them into audio files (.wav). We chose to include the full broadcast in extenso, encompassing both commentary simultaneous with sporting action and transition segments. The final corpus covers 18 curling events, from qualification rounds to finals, across men's, women's, and mixed competitions.

2.2 The formatting pipeline: from audio to CSV

The transcription of these audio files into textual data and, more generally, CSV files with metadata was one of the core challenges of our approach. To properly test our hypotheses regarding gendered biases, we also had to obtain information regarding the gender of both athletes and journalists.

The transcription was carried out using a Whisper model trained for the French language, enabling us to convert audio into timestamped text segments. The timestamping was crucial in order to articulate the obtained text with diarization. We chose to use a *medium* model as it offered a good trade-off between computing time and precision. The results are detailed, including punctuation, which is useful for making sentences our unit of analysis (cf section 3).

Speaker diarization was performed using pyannote.audio (pyannote/speaker-diarization-community-1), which identifies and separates individual speakers within each file. The natural tendency of pyannote to impute a large number of speakers that do not actually exist was not an issue in our case, though it could have been mitigated by constraining the model with a maximum number of speakers consistent with the type of content at hand.

As mentioned above, imputing the gender of speakers was crucial. To do so, we used InaSpeechSegmenter, an open-source tool developed by INA and David Doukhan. It classifies each audio segment

[1] https://www.dailymail.com/sport/olympics/article-13681573/Olympic-Games-commentator-REMOVED-Paris-coverage-sexist-remark.html

as *male*, *female*, *music*, *noise*, or *noEnergy*. The distribution of this classification across our dataset is described in subsection 2.4. Our pipeline includes a gender assignment rule of thumb based on majority overlap: the gender label with the highest cumulative overlap duration within the segment boundaries is selected as the main gender. A similar heuristic was applied to assign a unique speaker to each segment, as overlapping was observed but generally highly unbalanced.

The output is stored as one CSV per event, then concatenated into a single DataFrame with the following columns: segment start, segment stop, text, main speaker, speaker gender, and event ID. This ID enables us to merge the DataFrame with athlete gender metadata. We are therefore able to produce the descriptive statistics included in subsection 2.4.

2.3 Pre-processing

Once this single common CSV file obtained, we pre-processed the data in order to construct a solid basis for the upcoming analyses. This starts with filtering segments without spoken words: music, noise, and noEnergy segments identified by InaSpeechSegmenter are removed from the corpus (for the distribution, see Table 2).Only male and female speakers are retained. As we grouped text at the intervention level (cf. section 2), we divided it at the sentence level, using the sentence as the unit of analysis for topic modelling.

Compared to interventions, sentences are shorter and reduce the risk of introducing biases: each unit must be assigned a single topic, and longer interventions sometimes spanning several minutes might blur thematic boundaries and make unique assignment problematic.Words including capital letters are removed to limit as much as possible the inclusion of individual and country names that could bias embeddings. Unfortunately, Whisper occasionally writes proper nouns in lowercase. The consequences of which can be observed in the presence of names of country's inhabitants such as "suédois" (Topic 33 and 58) within some of the topics obtained (subsection 4.1, Table 10).

Finally, some interventions and therefore sentences were empty either originally or after the aforementioned cleaning steps, and are removed from the corpus.

2.4 Descriptive statistics

As mentioned, the total duration of broadcast audio is around 16 and a half hours across 18 events. Once converted to text, we obtain 12,486 segments, which constitutes a comprehensive dataset. The average segment length is 12 words with a standard deviation of 10.7. This is above the five-word threshold commonly used in the topic modelling literature (Li et al., 2020).

Event distribution (gender of athletes) There are 2 mixed events regarding the gender of athletes, 8 men's events, and the same number of women's events. Results in Table 1 show that in terms of commentary duration, men's events account for 59.6% (589 minutes) of the dataset while women's events account for 30.5% (301 minutes). Mixed events represent 9.9% (98 minutes). We observe a roughly 2:1 ratio in favour of men's events. This imbalance is not fundamentally an issue for our analysis, but should be kept in mind when interpreting subsequent results.

Table 1: Dataset distribution by athlete gender

	Men's (H)	**Women's (F)**	**Mixed (M)**	**Total**
Events	8	8	2	18
Duration (min)	589 (59.6%)	301 (30.5%)	98 (9.9%)	986
Sentences	7,201 (57.7%)	4,077 (32.7%)	1,208 (9.7%)	12,486
Avg. length (words)	12.1 (SD=11.1)	11.6 (SD=10.2)	12.1 (SD=10.0)	12.0 (SD=10.7)

Journalist gender distribution Regarding journalist gender (Table 2), male voices account for 64.7% of commentary duration (639 minutes) and 65.8% of segments (8,210), while female voices represent 31.2% (308 minutes) and 30.8% (3,844 segments). The remaining 4% corresponds to segments classified as music, noise, or noEnergy by InaSpeechSegmenter that are subsenquently exluded from the corpus (cf subsection 2.3). Female journalists tend to produce slightly longer sentences on average (13.3 words, SD=11.1) compared to male journalists (11.6 words, SD=10.5), though the difference remains modest.

Table 2: Dataset distribution by journalist gender

	Male	**Female**	**Other**	**Total**
Duration (min)	639 (64.7%)	308 (31.2%)	40 (4.1%)	986
Sentences	8,210 (65.8%)	3,844 (30.8%)	432 (3.5%)	12,486
Avg. length (words)	11.6 (SD=10.5)	13.3 (SD=11.1)	–	12.0 (SD=10.7)

Cross results: journalists' and athletes' gender Table 3 highlights a dominance of male journalists regardless of athlete gender. Remarkably, female journalists are quasi-absent from mixed events (12 segments vs. 1,185 for male). It is highly probable that both mixed events were in reality commented exclusively by male journalists, and that the 12 segments attributed to female voices are merely InaSpeechSegmenter misclassifications.

Table 3: Cross-tabulation of journalist gender and athlete gender

Table 4: *

Duration (minutes)

	F	**H**	**M**	**Total**
Female	85.1	221.7	0.9	307.7
Male	197.8	346.0	95.0	638.8
Music	12.9	9.1	0.1	22.1
noEnergy	1.3	2.7	0.6	4.6
Noise	3.4	9.1	0.9	13.4

Table 5: *

Segments

	F	**H**	**M**	**Total**
Female	1,069	2,763	12	3,844
Male	2,806	4,219	1,185	8,210
Music	135	88	1	224
noEnergy	18	35	6	59
Noise	49	96	4	149

Overall, this dataset is comprehensive though it exhibits a double imbalance: men's events receive approximately twice as much coverage as women's events, and male journalists account for roughly two thirds of all commentary. These structural asymmetries will be taken into account when interpreting the topic modelling results in the following sections.

3 Method: Choice of models

In regard to what we found in the literature, we designed an analytical approach enabling us to provide an original contribution. The main objective of our strategy is to obtain bottom-up exploratory results, as the literature mostly lacks recent works on our specific subject.

Topic Modeling We first conducted topic modelling to identify topics emerging from live commentary. We chose BERTopic as it exploits contextual embeddings rather than simple word co-occurrence as LDA or NMF do, allowing us to benefit from the full richness of our data. The main advantage of this approach is its capacity to capture the contextual meaning of words: a same word can carry different meanings depending on its context, and relying on such embeddings makes it possible to pursue a precise analysis of subtle and implicit latent sexist biases as we inteded to do.

Furthermore, BERTopic is particularly suitable for oral language corpora and relatively short sentences such as ours. Regarding the choice of embedding model, we use CamemBERT, a version of BERT pre-trained on a large French corpus, in order to maximise precision in capturing subtle and implicit meanings. More precisely, we use `dangvantuan/sentence-camembert-base`, a fine-tuned version designed for semantic similarity between sentences. We follow a standard

BERTopic pipeline: UMAP dimensionality reduction of embeddings, followed by HDBSCAN hierarchical density-based clustering, which avoids pre-supposing a number of topics and assigns ambiguous sentences to a catch-all noisy group (topic -1). Finally, c-TF-IDF helps identifies the most discriminating words for each cluster and therefore helps for its interpretation.

Each sentence being assigned to a single topic, we are able to observe the gendered distribution with respect to both athletes and journalists, and thus test potential disparities revealing underlying discursive biases.

Cosine similarity To further improve robustness, we complement this with a cosine similarity analysis. Topic distribution measures the prevalence of themes depending on gender, but two groups may use the same words — and therefore produce the same topics — in different semantic contexts without this affecting topic prevalence. Cosine similarity addresses this limitation by measuring the distance between embeddings, enabling us to test whether a same word is used in different semantic contexts depending on gender. This is particularly suited to our goals: we aim to identify subtle latent biases that may lie beyond simple lexical heterogeneity.

Robustness checks Finally, the robustness of our topic modelling is assessed via bootstrap analyses. As BERTopic is non-deterministic, topics may vary across runs on the same corpus. To assess stability, we sample 80% of the corpus 10 times and re-run the analysis on each sample. A topic is considered stable if it appears in at least 7 out of 10 runs with a Jaccard index above 0.2 on its keywords. This index measures the proportion of shared keywords between two topics across runs, serving as a proxy for topic consistency.

At this point, we had designed a topic modeling strategy that we estimate to be the most suitable for our goals: it is robustness-oriented and aims to provide substantive exploratory insights regarding gender biases in sport commentary. The results are detailed in section 4.

4 Results and discussion

In this section, I describe and interpret results obtained via the analytical strategy described in previous section 3.

4.1 Topic Modeling

General results I first applied the French version of BERTopic (CamemBERT) to my pre-processed dataset to obtain topics. As described in the dendrogram (Figure 1), the model produces 67 distinct topics plus one outlier topic. The latter contains 4,188 sentences, representing 35.1% of the dataset. This is substantial and may constitute a limitation of this analysis, potentially explained by the presence of shorter sentences and the highly technical vocabulary specific to curling. However, it should not undermine the rest of our analyses significantly.

The ten largest topics by size are presented in Table (Table 6). Topics 0, 4, and 7 consist mostly of technical curling vocabulary, which represents the core of live commentary. Topics 2 and 9 capture interjections and validation terms and do not provide meaningful insights for our purposes. Topic 8 focuses on the Olympic competition dimension through medal-related terms, which could be informative if its gendered distribution were significantly skewed. Topic 1 is noteworthy as it captures meta-discourse about curling itself, likely reflecting the explanatory role commentators play for a relatively unfamiliar sport. Finally, topics 12 and 13 (Figure 1) could prove meaningful in the presence of gender biases, as they reflect speakers' attention to match tension and closeness. Overall, the largest topics are dominated by technical vocabulary or generic oral language markers. As discussed below, the potentially most meaningful topics for our analysis are considerably smaller in size.

To assess the robustness of the topic structure, we ran BERTopic 10 times on 80% bootstrap samples of the corpus and measured topic stability using the Jaccard index. This procedure yields 28 stable topics out of 67. The stable topics cover the main themes identified in the analysis and especially technical curling vocabulary, competition, strategy, and nationality references. It thus reinforces the reliability of our interpretations while forcing us to be cautious while interpreting other topics.

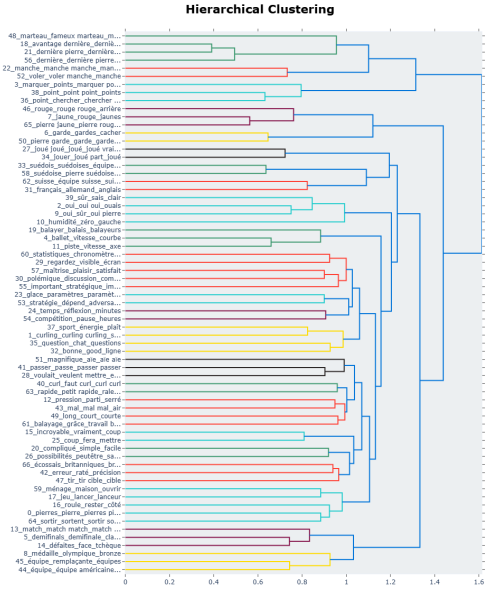

Figure 1: Dendogram of the topics obtained

Table 6: Ten largest topics produced by BERTopic

Topic	Keywords	Count
0	pierres, pierre, maison	458
1	curling, sport, classique	266
2	oui, ouais, informations	264
3	marquer, points, point	263
4	ballet, vitesse, courbe	254
5	demifinals, classement, équipes	239
6	garde, gardes, cacher	236
7	jaune, rouge, taper	236
8	médaille, olympique, bronze	223
9	oui, sûr, pierre	200
-1	*unclassified*	4,188

H1: Gender bias – Athletes As a reminder, only 32.7% of the sentences in our corpus were uttered while commenting women's events, compared to 57.7% for men's events. This distribution serves as a baseline to determine which topics are over- or under-represented for each athlete gender. It constitutes the first direct test of our first hypothesis (H1).

Table (Table 7) includes topics that show a significant deviation from this baseline (over 10% difference). Regarding topics overrepresented in women's events, topic 51 stands out as particularly interesting. It focuses on grace, aesthetics and emotion (*magnifique*, *superbe*) alongside exclamations (*aïe*), consistent with the literature on gendered sport commentary. BertTopic is not deterministic and therefore obtained topics may vary when we run it several times. It is interesting to note that, in other runs, the word "corps" (meaning body) was included in this topic suggesting an even more gendered

Table 7: Topics with notable gender deviation ($> | \pm 10\%|$) by athlete gender

Table 8: *

Overrepresented in women's events

Topic	Keywords	F_diff
48	marteau, fameux marteau	+21.6%
40	curl, faut curl	+19.2%
43	mal, mal mal, air	+15.1%
50	garde, protéger	+11.8%
20	compliqué, simple, facile	+11.5%
51	magnifique, aïe, superbe	+10.8%
24	temps, réflexion	+10.6%

Table 9: *

Overrepresented in men's events

Topic	Keywords	H_diff
66	écossais, britanniques	+30.4%
62	suisse, équipe suisse	+23.9%
27	joué joué, vraiment	+20.7%
14	défaites, gagner, face	+19.5%
53	stratégie, adversaire	+16.1%
18	avantage, dernière pierre	+14.1%
10	humidité, zéro	+15.1%

and even physiological dimension to this topic.

Topic 50 is also noteworthy, as it emphasizes defensive play (*garde*, *protéger* mean guard and protecting). Even if its offensive counterpart does not appear as clearly overrepresented in men's events, this gender gap could reflect latent gendered conceptions of female athletes as less aggressive.

Topic 48 (*marteau* meaning hammer) and topic 40 (*curl*) are harder to interpret as gender bias per se, as they may reflect specificities of certain events rather than systematic commentary differences.

Conversely, several topics overrepresented in men's events stand out. Topic 14 revolves around match results (*défaites*, *gagner* meaning loss and win), which could reflect greater journalistic investment in men's events and a greater focus on the outcome of games, though player nationality remains a potential confound.

Topic 27 groups superlative terms appearing more frequently in men's commentary.Topics 66 and 62, both related to player nationality (British and Swiss teams respectively), are more frequent in men's events but hardly constitute meaningful observations regarding gender bias. Finally, topic 53 (*stratégie*, *adversaire*) suggests that tactical commentary may be more prevalent in men's events.

To summarise, these first results are ambiguous and make it difficult to either confirm or reject H1 straightforwardly. One interpretation is that they reflect persisting latent gendered biases in sports commentary. Alternatively, this blurred picture may signal a further step in the trend toward gender equality in sport media coverage noted by Cooky et al., 2015. A further test is presented in Table 16 using cosine similarity to assess the robustness of these results.

H2: Gender bias – Journalists As for H1, the first step is to examine the baseline, namely the distribution of sentences according to commentator gender. This distribution mirrors the athletes' one: female journalists account for only 30.8% of segments, while male journalists cover 65.8%. As discussed below, the gender gap is larger for journalists than for athletes (cf Table 10), as more topics are significantly overrepresented with respect to commentator gender than athlete gender. This introduces a substantive nuance to previous findings in the sport commentary literature.

Regarding topics overrepresented among female commentators, it is notable that they over-invest technical, strategic, and descriptive dimensions of curling (topics 6, 7, 46, and 65). Their commentary may therefore be more rooted in objective and descriptive interventions.

Furthermore, topic 57 featuring words such as *plaisir* and *content* (pleasure, happy) suggests a potentially stronger emotional investment and expressiveness among female journalists, which would support H2.

Table 10: Topics with notable gender deviation ($> | \pm 10\%|$) by commentator gender

Table 11: *

Overrepresented among female commentators

Topic	Keywords	F_diff
6	garde, gardes, cacher	+32.9%
7	jaune, rouge, taper	+31.6%
53	maîtrise, plaisir, contente	+21.3%
46	rouge, arrière	+20.8%
18	avantage, dernière pierre	+18.0%
65	pierre jaune, couleur	+16.2%
63	rapide, ralenti, lent	+16.3%
57	maîtrise, plaisir, contente	+13.5%
45	équipe, féminine	+10.7%

Table 12: *

Overrepresented among male commentators

Topic	Keywords	H_diff
48	voler, manche, exploit	+28.4%
51	magnifique, aïe, superbe	+19.9%
56	dernière, manche dernière	+20.7%
35	question, chat	+17.1%
33	suédois, suédoises	+17.5%
58	suédoise, pierres suisses	+17.0%
66	écossais, britanniques	+18.8%
2	médaille, olympique	+15.4%
61	balayage, grâce, travail	+14.6%
22	dernière pierre, avantage	+14.6%

However, the overrepresentation of topic 51 (*magnifique*, *superbe*) among male commentators complicates this interpretation. This topic is nonetheless interesting: as previously mentioned, it is overrepresented in women's events regarding athlete gender, and in some bootstrap runs it featured the word *corps* (body). This could reflect marginal reminiscences of physiological commentary on female athletes by male journalists.

The picture for male journalists may be even more telling. Topic 61 (*balayage, grâce, travail*) is overrepresented among male commentators and in men's events, suggesting that the vocabulary of grace and precision is not reserved for women's commentary in this corpus. Topics 2 and 48 are perhaps the most substantive regarding the confirmation of H2: they illustrate a sharper focus on competition, athletic achievement, and performance among male journalists compared to their female counterparts, consistent with the existing literature detailed in the introduction.

First conclusion on H1 and H2 To summarise results for both H1 and H2, (Figure 2) simultaneously synthesises athletes and journalists' patterns, enabling a comparison of deviation sizes across topics. Consistently with our previous discussion, yellow bars (journalists) tend to be tendencially slightly longer than green ones (athletes), suggesting that commentator gender generates more substantial thematic differences than athlete gender. Deviations for athletes are moderate except for a few topics (51 and 48), and we therefore cannot confirm H1, consistently with the broader trend toward bias reduction in sport commentary.

Regarding journalists, while some topic deviations might seem surprising, we observe more significant and numerous deviations overall. This tends to provide slightly more support for H2 than for H1. Persisting biases are debatable but, if they exist in our data, appear to be concentrated on the journalists' side rather than the athletes' side.

These results are based on topic distributions alone. A complementary cosine similarity analysis is therefore desirable to test whether the same words are used in different contexts depending on the gender of athletes or journalists.

4.2 Cosine similarity: Robustness of our results

As mentioned, I produced two cosine similarity analyses: one for athletes and one for journalists. This design is made in order to test the robustness of the previous results. Cosine similarity enables us to assess whether the same words are used in different contexts depending on gender. It therefore provides a complementary test of H1 and H2.

Terms were selected in two ways: first, theoretically grounded words identified in the literature; second, the most frequent words in the corpus (Table 13). The second group ensures statistical robustness while the first directly tests topics identified in the literature.

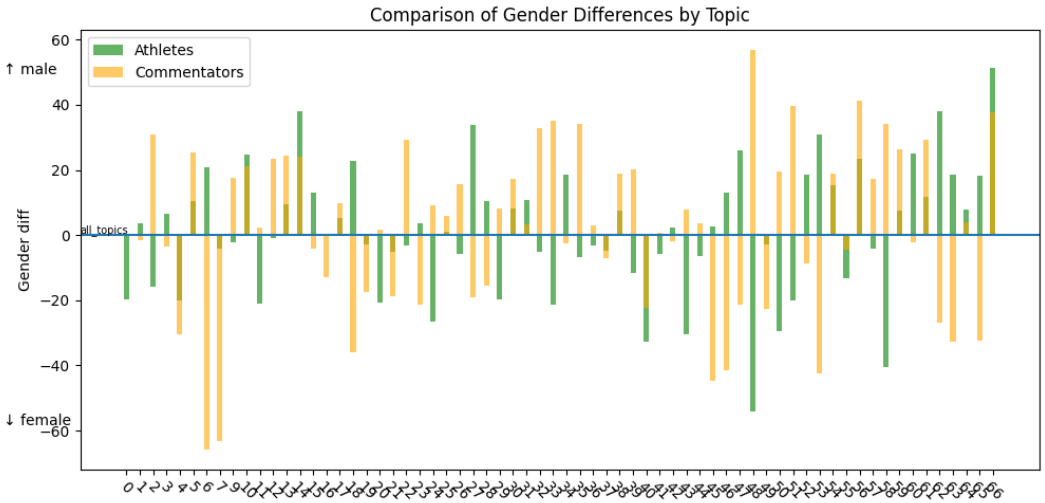

Figure 2: Comparison of gender differences by topic for athlete gender (green) and commentator gender (yellow). Positive values indicate overrepresentation in men's events or among male commentators; negative values indicate overrepresentation in women's events or among female commentators.

H1: Gender bias — Athletes For frequently occurring terms, cosine similarities are high (0.878 to 0.987, cf. Table 16), suggesting that this vocabulary is used in similar contexts regardless of athlete gender. This is consistent with the fact that these terms are predominantly technical. This step tends to confirm the absence of systematic bias and thus goes against H1.

Results are more nuanced for theoretically motivated terms. *Grâce* (grace) yields a similarity of 0.416 and *technique* 0.435 (Table 16). Both constitute meaningful signals suggesting these words are used in markedly different contexts depending on athlete gender. *Superbe* (0.510) and *effort* (0.560) show moderate differences.

However, all theoretically motivated terms are relatively rare in the corpus and do not appear in the top 50 most frequent words, making these estimates statistically less stable and calling for cautious interpretation.

H2: Gender bias — Journalists For frequently occurring terms, cosine similarities are slightly lower than for athletes (0.872 to 0.980, cf. Table 16). It suggests that commentator gender generates marginally more contextual differentiation than athlete gender for common vocabulary. *Maison* (0.872) and *capitaine* (0.893) show the most notable gaps, while technical curling terms such as *rouge*, *jaune* and *pierre* remain highly similar across genders.

Results are considerably more striking for theoretically motivated terms. *Grâce* yields a similarity of only 0.256 which is the lowest score across both analyses (athletes and journalists). It indicates that this word is used in radically different contexts depending on commentator gender. *Technique* (0.551), *effort* (0.679) and *superbe* (0.750) also show meaningful differences, stronger than their athlete-gender counterparts. This pattern tends to provide more support for H2 than the topic modelling results alone.

The same statistical limitation applies however: theoretically motivated terms are rare in the corpus and estimates should be interpreted with caution.

Taken together, both methods converge toward similar results as in subsection 4.1: the cosine similarity analysis reinforces the conclusion that H1 is not confirmed, while H2 finds slightly more empirical support. Crucially, the signal is concentrated on theoretically charged vocabulary, particularly *grâce*,

rather than on common technical terms, which is consistent with the latent and subtle nature of the biases described in the literature.

Table 13: Distribution of frequent terms by gender

Table 14: *

By athlete gender (baseline: M=58.4%, F=31.2%)

Term	**H%**	**F%**
capitaine	54	46
maison	36	64
vraiment	69	31
marquer	72	28
match	75	25
garde	65	35
rouge	66	34
jaune	69	31
pierre	60	40

Table 15: *

By commentator gender (baseline:M=64.5%,F=35.5%)

Term	**male%**	**female%**
maison	87	13
capitaine	55	45
garde	41	59
match	72	28
marquer	49	51
vraiment	55	45
rouge	39	61
jaune	36	64
pierre	64	36

Table 16: Cosine similarity between F and H centroids by term

Table 17: *

By athlete gender (baseline: M=58.4%, F=31.2%)

Term	**Cos. sim.**
Frequent terms	
capitaine	0.878
maison	0.924
vraiment	0.961
marquer	0.970
match	0.972
garde	0.973
rouge	0.976
jaune	0.983
pierre	0.987
Theoretical terms	
grâce	0.416
technique	0.435
superbe	0.510
effort	0.560
pression	0.843
stratégie	0.893
force	0.894
victoire	0.895
magnifique	0.951

Table 18: *

By commentator gender (baseline:M=64.5%,F=35.5%)

Term	**Cos. sim.**
Frequent terms	
maison	0.872
capitaine	0.893
garde	0.934
match	0.956
marquer	0.956
vraiment	0.967
rouge	0.972
jaune	0.977
pierre	0.980
Theoretical terms	
grâce	0.256
technique	0.551
effort	0.679
superbe	0.750
victoire	0.770
magnifique	0.830
stratégie	0.889
force	0.892
pression	0.894

5 Conclusion

This paper investigated the prevalence of latent sexist bias in French-language curling commentary during the 2026 Milan-Cortina Winter Olympic Games. My personal work is part of a broader collaborative effort covering multiple sports from the same broadcast corpus. I was specifically responsible for data retrieval, preprocessing, and descriptive statistics (cf section 5). However, once the code was made, each member of our group conducted its own topic modelling and cosine similarity analyses, made its own interpretation and redacted its own report.

Substantially, our results paint a nuanced picture. Regarding athlete gender (H1), thematic differences across topics are moderate and do not allow us to confirm the hypothesis of systematic gendered bias in how athletes are commented upon. This is broadly consistent with the trend toward reduced explicit and even latent sexism in sports media coverage identified by Cooky et al. (2015) and others. However, it is hard for me to be sure that this lack of discrepancy regarding athlete gender is caused more by substantial social trends than a lack of statistical power and therefore potentially merely artifactual.

Regarding commentator gender (H2), the picture is somewhat slightly different: more topics show significant deviations, and cosine similarity analyses reveal that theoretically charged terms are used in markedly different semantic contexts depending on whether the commentator is male or female. Therefore, if latent biases persist in our data, they appear to be concentrated on the journalists' side rather than being reflected in how athletes are framed.

Several limitations should be acknowledged. First, our topic modelling approach is fully unsupervised, which limits direct interpretability and makes it difficult to map topics onto specific theoretical themes identified in the literature. A supervised or semi-supervised approach with annotated data could strengthen the analysis. Second, we lack metadata on commentators beyond gender and information such as age, professional background, or whether a commentator is a journalist or a former athlete would allow for finer-grained analyses. Third, our results would benefit from cross-sport comparisons, controls for competition stage (qualification rounds vs. finals), and controls for the presence of French athletes, which may drive nationality-related commentary patterns independently of gender.

Future work could address these limitations by building a richer annotation framework, extending the corpus to other sports covered in this project, and integrating more granular metadata. Taken together, our findings suggest that the golden age of overt sexist commentary may indeed be over while remaining subtle, journalist-side discursive asymmetries appear to linger, warranting continued empirical attention.

Contributions

Our reports were jointly redacted, hence the potential similarities and common paragraphs. The interpretations of the results were nonetheless individual since we decided to apply our benchmark to a different sport each. We also worked together on the code, even though we had specific tasks attributed : Victor for data pre-processing and descriptive statistics, Lucas for topic modeling analysis and Quentin for semantic similarity analysis.

References

Billings, A. C., & Eastman, S. T. (2003). Framing identities: Gender, ethnic, and national parity in network announcing of the 2002 winter olympics. *Journal of Communication*. https://doi.org/10.1111/j.1460-2466.2003.tb02911.x

Capranica, L., & Aversa, F. (2002). Italian television sport coverage during the 2000 sydney olympic games: A gender perspective. *International Review for the Sociology of Sport*, *37*(3–4), 337–349.

Cooky, C., Messner, M., & Musto, M. (2015). "it's dude time!": A quarter century of excluding women's sports in televised news and highlight shows. *Communication & Sport*, *3*. https://doi.org/10.1177/2167479515588761

De Acutis, C., Weber, A., & Wurm, E. (2024). The effects of board gender quotas: A meta-analysis. *Labour Economics*, *91*, 102634. https://doi.org/10.1016/j.labeco.2024.102634

Li, L., Goh, T.-T., & Jin, D. (2020). How textual quality of online reviews affect classification performance: A case of deep learning sentiment analysis. *Neural Computing and Applications*, *32*, 4387–4415. https://doi.org/10.1007/s00521-019-04340-6

Messner, M., Duncan, M., & Jensen, K. (1993). Separating the men from the girls: The gendered language of televised sports. *Gender & Society*, *7*, 121–137. https://doi.org/10.1177/089124393007001007

Mutsaddi, A., Jamkhande, A., Thakre, A. S., & Haribhakta, Y. (2025). BERTopic for topic modeling of Hindi short texts: A comparative study. *Proceedings of the First Workshop on Natural Language Processing for Indo-Aryan and Dravidian Languages*, 22–32.

Rao, P., & Taboada, M. (2021). Gender bias in the news: A scalable topic modelling and visualization framework. *Frontiers in Artificial Intelligence*, *4*, 664737. https://doi.org/10.3389/frai.2021.664737

Zhou, K., Ethayarajh, K., Card, D., & Jurafsky, D. (2022). Problems with cosine as a measure of embedding similarity for high frequency words. https://arxiv.org/abs/2205.05092